

Limitations of the GradCafe Applicant Data Analysis

The GradCafe dataset consists of voluntarily submitted, anonymous applicant reports rather than a complete record of graduate-school applications and admission decisions. Applicants who choose to post their results may differ from those who do not, creating selection bias. For example, applicants with particularly positive or negative outcomes may be more motivated to submit a report. Individual applicants may also submit multiple entries, and the dataset does not necessarily identify or distinguish them. Consequently, the counts, averages, and acceptance percentages calculated from these records describe the submitted GradCafe entries, not the full population of graduate-school applicants. They should not be interpreted as official university admission rates or as reliable estimates of an individual applicant's probability of admission.

The dataset also depends on the accuracy and completeness of information entered by individual users. GPA and GRE scores may be missing, mistyped, or reported using inconsistent scales; university names, program names, degree types, application terms, and admission statuses may be entered in different formats. These inconsistencies can affect filtering, grouping, and comparisons across records. This analysis also uses LLM-generated fields to standardize or classify some applicant information. Although those fields make the data easier to query, the model may misinterpret ambiguous text or assign an incorrect classification. Differences between results based on original fields and LLM-generated fields illustrate why these classifications should not be treated as independently verified facts. The analysis therefore provides a useful summary of the available records, but its results remain subject to reporting errors, missing information, and classification uncertainty.